

## 第 4 章 Kafka API 实战

### 4.1 环境准备

- 1) 在 eclipse 中创建一个 java 工程
- 2) 在工程的根目录创建一个 lib 文件夹
- 3) 解压 kafka 安装包，将安装包 libs 目录下的 jar 包拷贝到工程的 lib 目录下，并 build path。
- 4) 启动 zk 和 kafka 集群，在 kafka 集群中打开一个消费者

```
[atguigu@hadoop102 kafka]$ bin/kafka-console-consumer.sh --zookeeper hadoop102:2181
--topic first
```

### 4.2 Kafka 生产者 Java API

#### 4.2.1 创建生产者（过时的 API）

```
package com.atguigu.kafka;
import java.util.Properties;
import kafka.javaapi.producer.Producer;
import kafka.producer.KeyedMessage;
import kafka.producer.ProducerConfig;

public class OldProducer {

    @SuppressWarnings("deprecation")
    public static void main(String[] args) {

        Properties properties = new Properties();
        properties.put("metadata.broker.list", "hadoop102:9092");
        properties.put("request.required.acks", "1");
        properties.put("serializer.class", "kafka.serializer.StringEncoder");

        Producer<Integer, String> producer = new Producer<Integer,String>(new
        ProducerConfig(properties));

        KeyedMessage<Integer, String> message = new KeyedMessage<Integer,
        String>("first", "hello world");
        producer.send(message);
    }
}
```

## 4.2.2 创建生产者（新 API）

```
package com.atguigu.kafka;
import java.util.Properties;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.Producer;
import org.apache.kafka.clients.producer.ProducerRecord;

public class NewProducer {

    public static void main(String[] args) {

        Properties props = new Properties();
        // Kafka 服务端的主机名和端口号
        props.put("bootstrap.servers", "hadoop103:9092");
        // 等待所有副本节点的应答
        props.put("acks", "all");
        // 消息发送最大尝试次数
        props.put("retries", 0);
        // 一批消息处理大小
        props.put("batch.size", 16384);
        // 请求延时
        props.put("linger.ms", 1);
        // 发送缓存区内存大小
        props.put("buffer.memory", 33554432);
        // key 序列化
        props.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
        // value 序列化
        props.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");

        Producer<String, String> producer = new KafkaProducer<>(props);
        for (int i = 0; i < 50; i++) {
            producer.send(new ProducerRecord<String, String>("first", Integer.toString(i),
"hello world-" + i));
        }

        producer.close();
    }
}
```

### 4.2.3 创建生产者带回调函数（新 API）

```
package com.atguigu.kafka;
import java.util.Properties;
import org.apache.kafka.clients.producer.Callback;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.clients.producer.RecordMetadata;

public class CallBackProducer {

    public static void main(String[] args) {

        Properties props = new Properties();
        // Kafka 服务端的主机名和端口号
        props.put("bootstrap.servers", "hadoop103:9092");
        // 等待所有副本节点的应答
        props.put("acks", "all");
        // 消息发送最大尝试次数
        props.put("retries", 0);
        // 一批消息处理大小
        props.put("batch.size", 16384);
        // 增加服务端请求延时
        props.put("linger.ms", 1);
        // 发送缓存区内存大小
        props.put("buffer.memory", 33554432);
        // key 序列化
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        // value 序列化
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

        KafkaProducer<String, String> kafkaProducer = new KafkaProducer<>(props);

        for (int i = 0; i < 50; i++) {

            kafkaProducer.send(new ProducerRecord<String, String>("first", "hello" + i), new
Callback() {

                @Override
                public void onCompletion(RecordMetadata metadata, Exception exception) {

                    if (metadata != null) {

                        System.err.println(metadata.partition() + "---" + metadata.offset());
```

```
        }  
    }  
    });  
}  
  
kafkaProducer.close();  
}  
}
```

#### 4.2.4 自定义分区生产者

0) 需求：将所有数据存储到 topic 的第 0 号分区上

1) 定义一个类实现 `Partitioner` 接口，重写里面的方法（过时 API）

```
package com.atguigu.kafka;  
import java.util.Map;  
import kafka.producer.Partitioner;  
  
public class CustomPartitioner implements Partitioner {  
  
    public CustomPartitioner() {  
        super();  
    }  
  
    @Override  
    public int partition(Object key, int numPartitions) {  
        // 控制分区  
        return 0;  
    }  
}
```

2) 自定义分区（新 API）

```
package com.atguigu.kafka;  
import java.util.Map;  
import org.apache.kafka.clients.producer.Partitioner;  
import org.apache.kafka.common.Cluster;  
  
public class CustomPartitioner implements Partitioner {  
  
    @Override  
    public void configure(Map<String, ?> configs) {  
  
    }  
  
    @Override
```

```
public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[] valueBytes, Cluster cluster) {  
    // 控制分区  
    return 0;  
}  
  
@Override  
public void close() {  
  
}  
}
```

### 3) 在代码中调用

```
package com.atguigu.kafka;  
import java.util.Properties;  
import org.apache.kafka.clients.producer.KafkaProducer;  
import org.apache.kafka.clients.producer.Producer;  
import org.apache.kafka.clients.producer.ProducerRecord;  
  
public class PartitionerProducer {  
  
    public static void main(String[] args) {  
  
        Properties props = new Properties();  
        // Kafka 服务端的主机名和端口号  
        props.put("bootstrap.servers", "hadoop103:9092");  
        // 等待所有副本节点的应答  
        props.put("acks", "all");  
        // 消息发送最大尝试次数  
        props.put("retries", 0);  
        // 一批消息处理大小  
        props.put("batch.size", 16384);  
        // 增加服务端请求延时  
        props.put("linger.ms", 1);  
        // 发送缓存区内存大小  
        props.put("buffer.memory", 33554432);  
        // key 序列化  
        props.put("key.serializer",  
"org.apache.kafka.common.serialization.StringSerializer");  
        // value 序列化  
        props.put("value.serializer",  
"org.apache.kafka.common.serialization.StringSerializer");  
        // 自定义分区  
        props.put("partitioner.class", "com.atguigu.kafka.CustomPartitioner");  
    }  
}
```

```
Producer<String, String> producer = new KafkaProducer<>(props);
producer.send(new ProducerRecord<String, String>("first", "1", "atguigu"));

producer.close();
}
}
```

#### 4) 测试

(1) 在 hadoop102 上监控/opt/module/kafka/logs/目录下 first 主题 3 个分区的 log 日志动态变化情况

```
[atguigu@hadoop102 first-0]$ tail -f 00000000000000000000.log
```

```
[atguigu@hadoop102 first-1]$ tail -f 00000000000000000000.log
```

```
[atguigu@hadoop102 first-2]$ tail -f 00000000000000000000.log
```

(2) 发现数据都存储到指定的分区了。

## 4.3 Kafka 消费者 Java API

#### 0) 在控制台创建发送者

```
[atguigu@hadoop104 kafka]$ bin/kafka-console-producer.sh --broker-list hadoop102:9092
--topic first
>hello world
```

#### 1) 创建消费者 (过时 API)

```
package com.atguigu.kafka.consume;
import java.util.HashMap;
import java.util.List;
import java.util.Map;
import java.util.Properties;
import kafka.consumer.Consumer;
import kafka.consumer.ConsumerConfig;
import kafka.consumer.ConsumerIterator;
import kafka.consumer.KafkaStream;
import kafka.javaapi.consumer.ConsumerConnector;

public class CustomConsumer {

    @SuppressWarnings("deprecation")
    public static void main(String[] args) {
        Properties properties = new Properties();
```

```
properties.put("zookeeper.connect", "hadoop102:2181");
properties.put("group.id", "g1");
properties.put("zookeeper.session.timeout.ms", "500");
properties.put("zookeeper.sync.time.ms", "250");
properties.put("auto.commit.interval.ms", "1000");

// 创建消费者连接器
ConsumerConnector consumer = Consumer.createJavaConsumerConnector(new
ConsumerConfig(properties));

HashMap<String, Integer> topicCount = new HashMap<>();
topicCount.put("first", 1);

Map<String, List<KafkaStream<byte[], byte[]>>> consumerMap =
consumer.createMessageStreams(topicCount);

KafkaStream<byte[], byte[]> stream = consumerMap.get("first").get(0);

ConsumerIterator<byte[], byte[]> it = stream.iterator();

while (it.hasNext()) {
    System.out.println(new String(it.next().message()));
}
}
```

## 2) 官方提供案例（自动维护消费情况）（新 API）

```
package com.atguigu.kafka.consume;
import java.util.Arrays;
import java.util.Properties;
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.KafkaConsumer;

public class CustomNewConsumer {

    public static void main(String[] args) {

        Properties props = new Properties();
        // 定义 kafka 服务的地址，不需要将所有 broker 指定上
        props.put("bootstrap.servers", "hadoop102:9092");
        // 制定 consumer group
        props.put("group.id", "test");
        // 是否自动确认 offset
```

```
        props.put("enable.auto.commit", "true");
        // 自动确认 offset 的时间间隔
        props.put("auto.commit.interval.ms", "1000");
        // key 的序列化类
        props.put("key.deserializer",
"org.apache.kafka.common.serialization.StringDeserializer");
        // value 的序列化类
        props.put("value.deserializer",
"org.apache.kafka.common.serialization.StringDeserializer");
        // 定义 consumer
        KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);

        // 消费者订阅的 topic, 可同时订阅多个
        consumer.subscribe(Arrays.asList("first", "second", "third"));

        while (true) {
            // 读取数据, 读取超时时间为 100ms
            ConsumerRecords<String, String> records = consumer.poll(100);

            for (ConsumerRecord<String, String> record : records)
                System.out.printf("offset = %d, key = %s, value = %s%n", record.offset(),
record.key(), record.value());
        }
    }
```